

Project2

April 30, 2022

```
[1]: import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import geopandas as gpd
import scipy.stats
from scipy.stats import norm
```

I chose this topic because I am a runner. I haven't run in many years but recently got the itch to run again. I wanted to do something that I hadn't done before. The longest race I have ever done is 25 kilometers or about 16.6 miles. It seemed only natural to take on a marathon. I chose the San Francisco Marathon in July. I tracked my exercise numbers and we will look at historical data to find some trends.

The first dataset is a dataset that I created in February to track my daily exercise numbers. The two other datasets we will look at are from the Berlin marathon from 1974 to 2019. The first dataset has individual runner data and the second has yearly weather data from the race day.

V, David. (2019). Berlin marathon data. [data files] www.bmw-berlin-marathon.com
<https://www.kaggle.com/datasets/aiaiaidavid/berlin-marathons-data/>

```
[4]: marathon = pd.read_csv("C:\\Users\\Brett\\Desktop\\ASU\\Classes\\Spring_
↳2022\\DAT 301\\Project2\\Berlin_Marathon_data_1974_2019.csv",
↳low_memory=False)
weather = pd.read_csv("C:\\Users\\Brett\\Desktop\\ASU\\Classes\\Spring_
↳2022\\DAT 301\\Project2\\Berlin_Marathon_weather_data_since_1974.csv",
↳low_memory=False)
mylog = pd.read_csv("C:\\Users\\Brett\\Desktop\\ASU\\Classes\\Spring 2022\\DAT_
↳301\\Project2\\mylog.csv", low_memory=False)
marathon = marathon.loc[marathon["TIME"] != 'no time']
```

This is my exercise training data

```
[5]: display(mylog)
```


	DATE	RUN/Miles	WALK/Miles	BIKE/Miles	SWIM/Minutes	WEIGHT/Pounds
0	2022-2-11	2.0	0.0	0	0	208
1	2022-2-12	3.0	0.0	0	0	206
2	2022-2-13	3.0	0.0	0	0	206
3	2022-2-14	3.0	0.0	0	0	205

4	2022-2-15	2.0	0.0	0	0	205
..	
74	2022-4-26	8.5	0.5	0	0	167
75	2022-4-27	7.0	0.5	0	0	166
76	2022-4-28	8.5	0.5	0	0	166
77	2022-4-29	11.0	0.5	0	0	167
78	2022-4-30	8.5	0.5	0	0	165

[79 rows x 6 columns]

```
[20]: fig, (ax1, ax2) = plt.subplots(2, 1, figsize = (12,12))
mylog.plot(x = 'DATE', y = ['RUN/Miles', 'WALK/Miles', 'BIKE/Miles'], ax=ax1)
ax1.set_title("Daily mileage totals during training")
ax1.set_ylabel("MILES")
mylog.plot(x = 'DATE', y = ['WEIGHT/Pounds'], ax=ax2, color = 'red')
ax2.set_title("Daily weight loss during training")
ax2.set_ylabel("POUNDS")
```

```
[20]: Text(0, 0.5, 'POUNDS')
```



At first, the training was bogged down by tendonitis. The big spike in biking miles was a reaction to this. I tried hiking in March that made me horribly tired. Luckily my running form came back and I was able to resume running. I was also able to resume biking. In April, the biking was getting boring and wasn't that useful for running anymore so I ditched it. Training for the marathon is going well now!

The dataset below is from the Berlin Marathon. It has been merged and cleaned. It also has the finishing times in convenient units for measurement.

```
[16]: def fun(time):
      return int(time[0:2]) * 3600 + int(time[3:5]) * 60 + int(time[6:8])

Time_in_seconds = marathon['TIME'].apply(fun)
marathon['Time_in_seconds'] = Time_in_seconds
```

```

marathon['Time_in_hours'] = Time_in_seconds / 3600
test = marathon.groupby('YEAR', as_index=False)['Time_in_hours'].mean()
test.rename(columns = {'Time_in_hours': 'Avg_time_hours'}, inplace = True)
marathon_merged = pd.merge(marathon, test)
marathon_merged = pd.merge(marathon_merged, weather)
marathon_merged = marathon_merged.sort_values(by = ['Time_in_hours'])
display(marathon_merged)

```

	YEAR	COUNTRY	GENDER	AGE	TIME	Time_in_seconds	Time_in_hours	\
811102	2018	NaN	male	30	02:01:39	7299	2.027500	
851743	2019	ETH	male	35	02:01:41	7301	2.028056	
851744	2019	ETH	male	20	02:02:48	7368	2.046667	
670190	2014	NaN	male	30	02:02:57	7377	2.049167	
735888	2016	NaN	male	30	02:03:03	7383	2.050833	
...	
771878	2016	NaN	male	60	08:41:33	31293	8.692500	
811101	2017	NaN	male	60	08:47:19	31639	8.788611	
216924	1998	NaN	female	55	09:19:54	33594	9.331667	
216925	1998	NaN	female	60	09:43:23	35003	9.723056	
216926	1998	NaN	female	40	09:49:41	35381	9.828056	

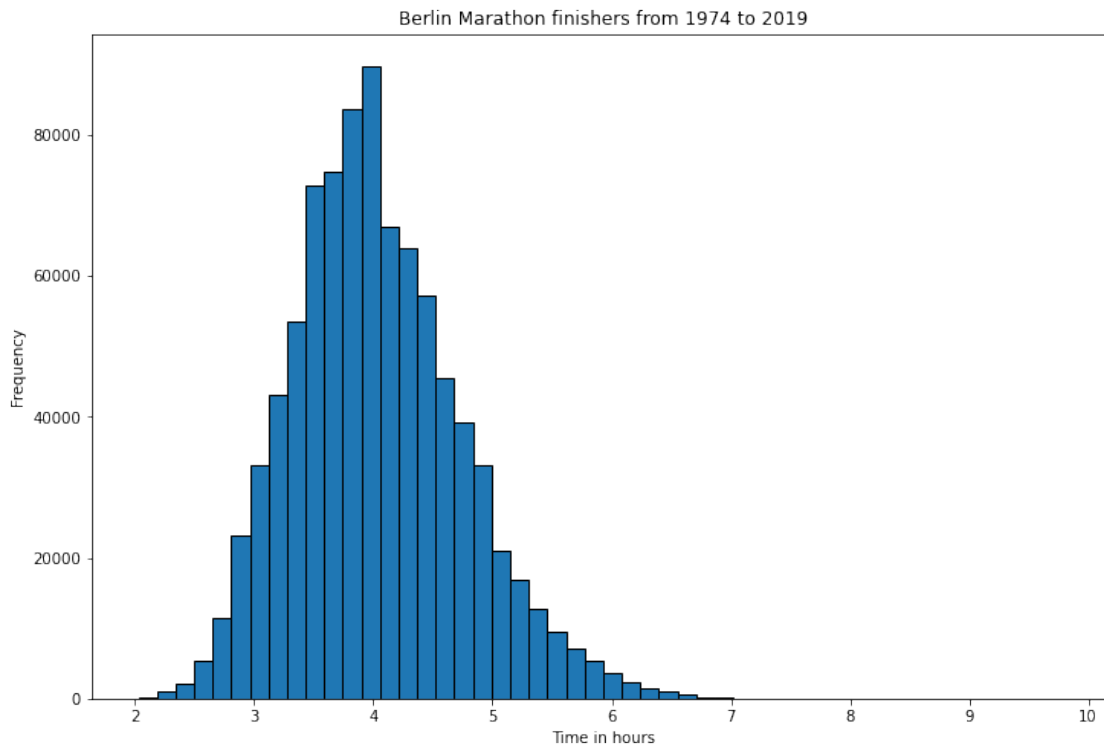
	Avg_time_hours	PRECIP_mm	SUNSHINE_hrs	CLOUD_hrs	ATMOS_PRESS_mbar	\
811102	4.316632	0.0	7.75	5.6	1013.82	
851743	4.044367	8.0	0.70	6.8	993.62	
851744	4.044367	8.0	0.70	6.8	993.62	
670190	4.130463	0.0	9.78	4.5	1016.39	
735888	4.210615	0.0	10.43	5.6	1012.05	
...	
771878	4.210615	0.0	10.43	5.6	1012.05	
811101	4.204538	0.3	1.50	7.5	1012.53	
216924	3.961136	0.0	9.90	1.0	1026.80	
216925	3.961136	0.0	9.90	1.0	1026.80	
216926	3.961136	0.0	9.90	1.0	1026.80	

	AVG_TEMP_C	MAX_TEMP_C	MIN_TEMP_C
811102	14.9	22.4	9.0
851743	14.0	16.4	11.2
851744	14.0	16.4	11.2
670190	12.7	21.2	5.1
735888	14.5	23.0	8.3
...
771878	14.5	23.0	8.3
811101	14.2	16.4	12.3
216924	13.8	19.8	10.3
216925	13.8	19.8	10.3
216926	13.8	19.8	10.3

[882539 rows x 15 columns]

```
[17]: fig, ax = plt.subplots(1, 1, figsize = (12,8))
marathon_merged = marathon_merged.sort_values(by = ['YEAR'])
ax.hist(marathon_merged['Time_in_hours'], histtype='bar', bins = 50,
        edgecolor='black', linewidth=1)
ax.set_title("Berlin Marathon finishers from 1974 to 2019")
ax.set_xlabel("Time in hours")
ax.set_ylabel("Frequency")
```

```
[17]: Text(0, 0.5, 'Frequency')
```




The graph above shows all the finishers of the Berlin Marathon over the years. You can clearly see the mean finishing time is about 4 hours. My goal is to do better than that.

```
[19]: marathon_merged_male = marathon_merged[marathon_merged['GENDER'] == 'male']
marathon_merged_female = marathon_merged[marathon_merged['GENDER'] == 'female']
fig, ax = plt.subplots(1, 1, figsize = (12,8))
ax.hist(marathon_merged_male['Time_in_hours'], histtype='bar', bins = 50,
        edgecolor='black', linewidth=1, label='male', color = 'blue')
ax.hist(marathon_merged_female['Time_in_hours'], histtype='bar', bins = 50,
        edgecolor='black', linewidth=1, label='female', color = 'red')
ax.legend()
```

```
ax.set_title("Berlin Marathon finishers from 1974 to 2019 by Gender")
ax.set_xlabel("Time in hours")
ax.set_ylabel("Frequency")
```


```
[19]: Text(0, 0.5, 'Frequency')
```



There is a clear difference here in speed and in the number of competitors between men and women.

```
[5]: weather.plot(x = 'YEAR', y = ['PRECIP_mm', 'AVG_TEMP_C'], figsize=(12, 8),
    ↳ ylabel = "Celcius or mm of rain")
plt.title("Avg precipitation and avg temperature for Berlin marathoners for
    ↳ each year")
```


```
[5]: Text(0.5, 1.0, 'Avg precipitation and avg temperature for Berlin marathoners for
    each year')
```



From this graph we can see the general numbers for precipitation and avg temperatures during each years marathon in Berlin.

```
[6]: fig, ax = plt.subplots(1, 1, figsize = (12,8))
marathon_merged.plot.scatter(x= 'YEAR' , y='Time_in_hours', c='AVG_TEMP_C',
    ↪ax=ax, colormap='turbo', ylim = [2,7])
marathon_merged.plot(x = 'YEAR', y = 'Avg_time_hours', ax=ax, color = 'black',
    ↪linewidth = 3, linestyle='dashed')
plt.title("Avg time in hours for Berlin marathoners for each year against the
    ↪average temperature")
```

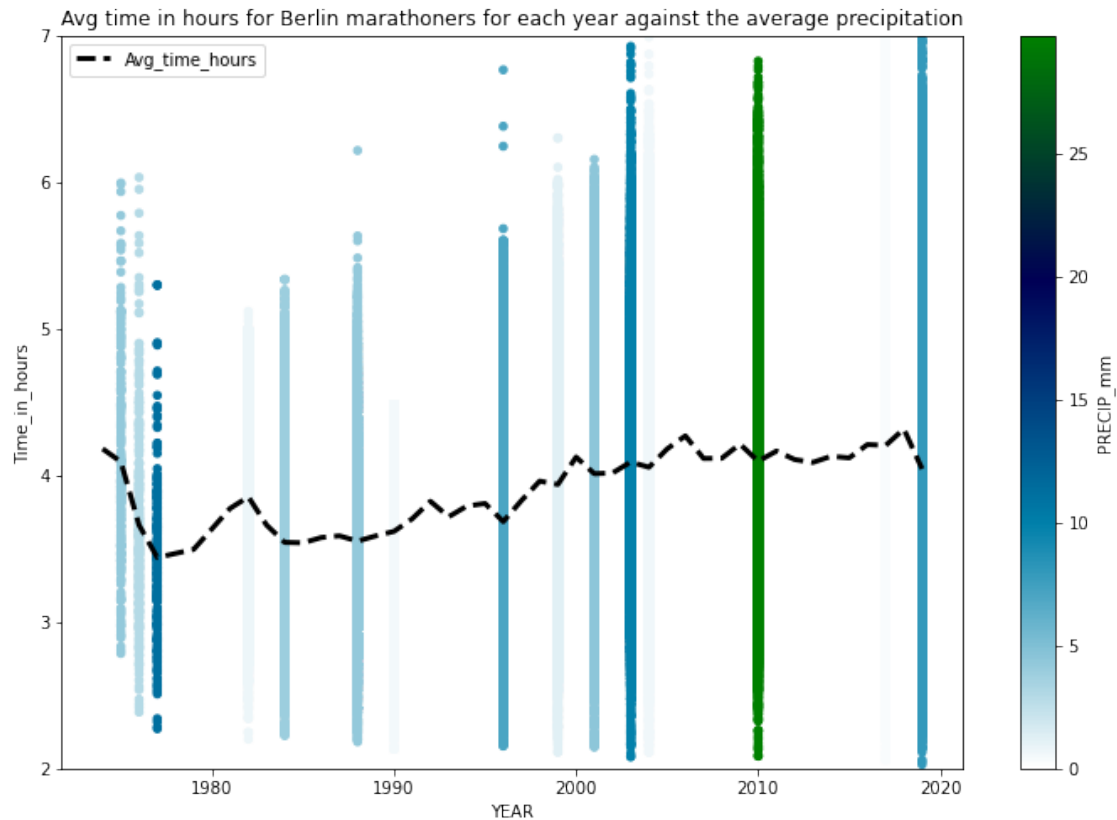
```
[6]: Text(0.5, 1.0, 'Avg time in hours for Berlin marathoners for each year against
the average temperature')
```



From this graph we can make some interesting conclusions. First, we can see a strong correlation between the temperature and the average finishing time (the dashed line in black). As the temperature increases we see an increase in the finishing time. We also see a large increase in finishing time when the temperature is extremely cold. Mild decreases in temperature seem to yield better marathon finishing time results.

```
[21]: fig, ax = plt.subplots(1, 1, figsize = (12,8))
marathon_merged.plot.scatter(x= 'YEAR' , y='Time_in_hours', c='PRECIP_mm',
    ↪ax=ax, colormap='ocean_r', ylim = [2,7])
marathon_merged.plot(x = 'YEAR', y = 'Avg_time_hours', ax=ax, color = 'black',
    ↪linewidth = 3, linestyle='dashed')
plt.title("Avg time in hours for Berlin marathoners for each year against the
    ↪average precipitation")
```

```
[21]: Text(0.5, 1.0, 'Avg time in hours for Berlin marathoners for each year against
the average precipitation')
```

Maybe unexpectedly, precipitation seems to have an even stronger correlation with the finishing times. When it rains, the average finishing times decrease the most. Even in years where the average temperature is warm, the average finishing time lowers when there is a large amount of precipitation. This might be caused by the water acting as a cooling agent, allowing the body to cool itself without having to sweat, which saves energy and allows the runners to run faster.

In conclusion, running a marathon is a challenge. Most runners take about 4 hours to complete one. Men run a bit faster. Also, cooler course temperature and sufficient rainfall have about a 15 to 30 minute effect on average finishing times. San Francisco will be very cool as the average morning temperature in July is about 55 degrees fahrenheit. This is almost exactly the same as Berlin's average race temperature. However, San Francisco is a much hillier course and I will have to take that into account.

[]: